Robust Subgraph Generation Improves Abstract Meaning Representation Parsing

Abstract

Abstract Meaning Representation (AMR) is a representation for open-domain rich semantics, with potential use in fields like semantic parsing and machine translation. We identify that node generation, typically done using a simple dictionary lookup, is currently the crucial limiting factor in AMR parsing. We propose a small set of actions to construct parts of the AMR parse from spans of tokens, which allows for more robust learning of this stage. We show that our set of construction actions generalize better than the previous approach, even when learned with an extremely simple classifier. We improve on the previous state-of-the-art result for AMR parsing, boosting end-to-end F1 from 0.59 to 0.62 on the LDC2013E117 and LDC2014T12 datasets.

1 Introduction

Abstract Meaning Representation (AMR) (Banarescu et al., 2013) is a rich graph-based language for expressing semantics over a broad domain. Figure 1 shows an example AMR for "he gleefully ran to his dog Rover", and we give a brief tutorial on AMR in Section 2.

AMR parsing is the task of mapping a natural language sentence into an AMR graph. AMR parsing is exciting because a practical broad-domain AMR parser could enable a new breed of natural language applications ranging from semantically aware MT to rich broad-domain QA over text-based knowledge bases. At the time of this writing AMR is the target of a multi-institution multi-year data-labeling program led by Kevin Knight, which promises to produce a corpus that is large enough to make it possible to parse to AMR well despite its many challenges.


Figure 1: The AMR graph for "He gleefully ran to his dog Rover". Nodes represent concepts, and arcs are relationships between concepts. The dark arc labeled "op1" is expected to be generated by NER++.

However, no matter how much data we are soon to have, AMR parsing remains a hard unsolved task, with only modest performance numbers reported so far. A good AMR parser needs to be able to handle a broad range of semantic interpretation tasks, and AMR also possesses distinctive computational challenges from structural properties such as providing no guarantees about projectivity or even acyclicity of its graphs.

We follow previous work (Flanigan et al., 2014) in dividing AMR parsing into two steps. The first step is *concept identification*, which generates concept (or entity) nodes from text, and which we'll refer to as *NER++* (Section 3.1). The sec-
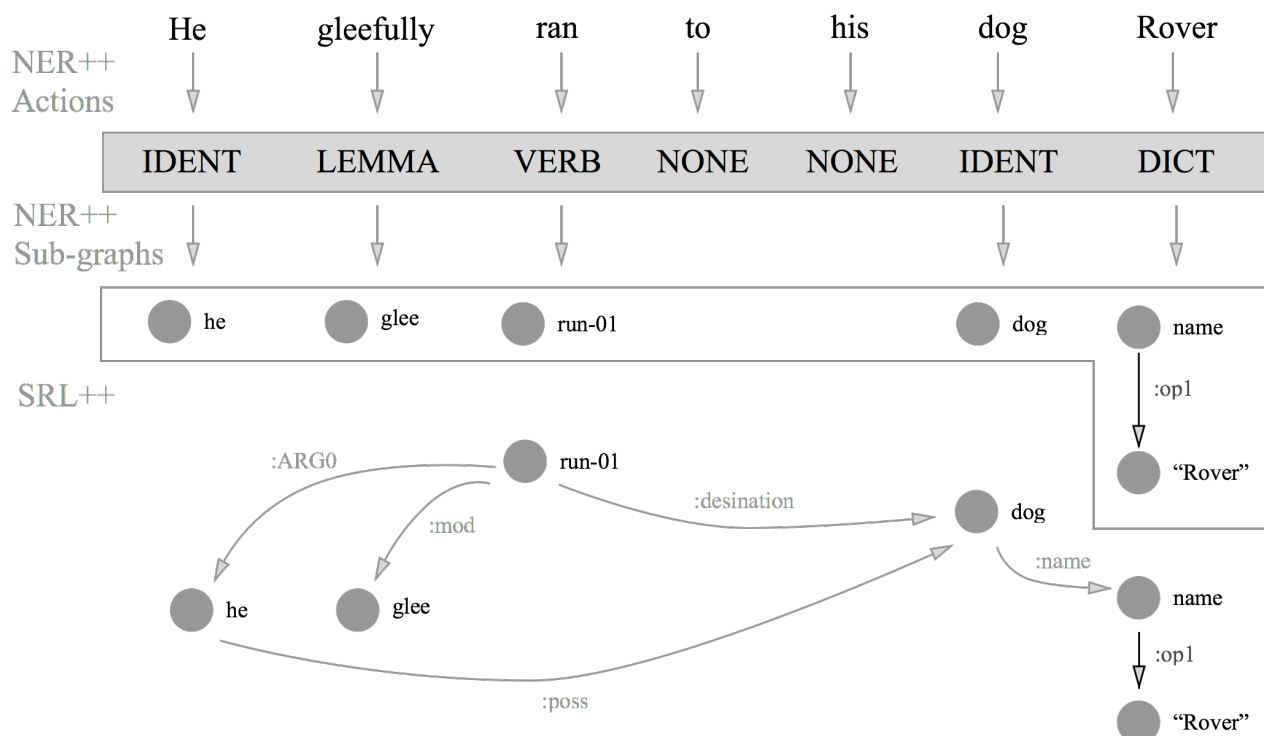
Figure 2: Derivation process for "He gleefully ran to his dog Rover". First the tokens in the sentence are labeled with derivation actions, then those actions are used to generate AMR sub-graphs, and then those sub-graphs are stitched together to form a coherent whole.

ond step is *relation identification*, which adds arcs to link these nodes into a fully connected AMR graph, which we'll call *SRL++* (Section 3.2).

Our initial assumption (which we suspect would be shared by many researchers with a background in structured prediction) was to assume that the main area deserving further attention in AMR parsing is predicting the edges in the not well-aligned, cyclic non-projective graphs of AMR (i.e., SRL++). This intuition follows by analogy from the locus of work for syntactic dependency parsing. We spent months testing novel structure-prediction algorithms and training regimes to capture aspects of the non-projective cyclic graphs, only to find that the careful application of MST presented in Flanigan et al. (2014) set an extremely strong baseline for the SRL++ task, and our gains were negligible to non-existent.

Over time we realized that SRL++ is *not* the hard part of AMR parsing. The hard part of AMR parsing is NER++. A first piece of evidence for this claim comes from a close analysis of the only published AMR parser to date, (Flanigan et al., 2014), dubbed "JAMR". When JAMR parses using its own NER++ and SRL++ components,

it gets an end-to-end score of 0.58 F1. However, when JAMR is given a gold NER++ output, and must only perform SRL++ over given sub-graphs it scores 0.80 F1 against a virtual upper bound of 0.83 F1, which is inter-annotator agreement on AMR **TODO: NEED CITE FOR THIS** . (And annotators very rarely disagree on NER++, which amounts to disagreeing about the entities mentioned in a sentence). **TODO:** verify – CDM: I'm not sure this is true [since surely there is quite a bit of inconsistency in NER++, such as for which compound words with derivational morphology to expand into subgraphs.] So this means the SRL++ component is nearly perfect, if given perfect NER++.

In retrospect, it makes sense that SRL++ within AMR is *relatively* easy given a perfect set of nodes to link together. There's a strong type-check feature for the existence and type of any arc just by looking at its end-points, and syntactic dependency features are very informative for removing any remaining ambiguity. If a system is considering how to link the node "run-01" in Figure 1, the verb-sense frame for "run-01" leaves very little entropy in terms of what we could assign as an

ARG0 arc. It must be a noun, which leaves either "he" or "dog", and this is easily decided in favor of "he" by looking for an nsubj arc in the dependency parse. Given a modest amount of data, a simple parser can learn these lexical and syntactic type-check rules.

While SRL++ is hard, NER++ is extremely challenging, encompassing many SemEval tasks in a single monolith. This task was handled by a simple baseline system in (Flanigan et al., 2014), which memorized a purely lexical mapping from spans of text to AMR nodes, augmented by an NER system and time expression regex at test time. This is problematic, because fully 38% of the tokens in the test set are unobserved at training time, and the existing method has no way to handle generation for them.

The primary contribution of this paper is a method to largely delexicalize NER++ to gracefully handle unseen tokens. Choosing from among a small set of 'generative actions' our system can derive an AMR sub-graph from a span of tokens (see Figure 2). For example, we have an action *VERB* that will perform a verb-sense-disambiguation to the appropriate PropBank frame (**?**)n the source token, like the "ran" to "run-01" example in Figure 1.

We show that this approach improves recall dramatically over previous approaches, and that end to end performance is improved from 0.59 to 0.62 smatch when our generative actions are stitched together by the previous state of the art parser (Flanigan et al., 2014). We suspect that the gains are not even bigger mainly because our NER++ system is able to generate nodes that the SRL++ system has never seen during training time, which makes the still lexicalized typecheck features of the SRL++ component useless.

2 A Crash-Course in AMR

AMR is a language for expressing semantics as a rooted, directed, and potentially cyclic graph, where nodes represent concepts and arcs are relationships between concepts. The nodes (concepts) in an AMR graph do not have to be explicitly grounded in the source sentence, and while such an alignment can often be generated it is not provided in the training corpora. The semantics of nodes can represent lexical items (e.g., *dog*), sense tagged lexical items (e.g., *run-01*), type markers (e.g., *date-entity*), and a host of other phenomena.

The edges (relationships) in AMR describe one of a number of semantic relationships between concepts. The most salient of these is semantic role labels, such as the ARG0 and destination arcs in Figure 2. However, often these arcs define a semantics more akin to syntactic dependencies (e.g., *mod* standing in for adjective and adverbial modification), or take on domain-specific semantics (e.g., the month, day, and year arcs of a *date-entity*).

AMR is based on neo-Davidsonian semantics, (Davidson, 1967; Parsons, 1990). To introduce AMR and its notation in more detail, we'll unpack the translation of the sentence "he gleefully ran to his dog Rover". We show in Figure 1 the interpretation of this sentence as an AMR graph.

The root node of the graph is labeled *run-01*, corresponding to the PropBank (**?**) definition of the verb *ran* . *run-01* has an outgoing *ARG0* arc to a node *he*, with the usual PropBank semantics. The outgoing *mod* edge from *run-01* to *glee* takes a general purpose semantics corresponding to adjective, adverbial, or other modification of the governor by the dependent. We note that *run-01* has a *destination* arc to *dog*. The label for *destination* is taken from a finite set of special arc sense tags similar to the preposition senses found in (**?**). The last portion of the figure parses *dog* to a node which will eventually serve as a type marker, and *Rover* into the larger subgraph indicating a concept with name "Rover."

2.1 Formal task definition

Formally, an AMR graph is represented as a directed multigraph. Following Flanigan et al. (2014), we simplify the representation by (1) disallowing cycles in the graph, and (2) disallowing multiple edges between two nodes in a graph. Both of these are rare phenomena in the representation, and substantially complicate the task.

Thus, we treat AMR as a directed graph G of nodes $\mathbf{N} = \{n_0 \ldots n_k\}$, and an edge matrix $\mathbf{A}$ such that $\mathbf{A}_{i,j} = l$ means that an arc exists from n_i to n_j with label l; $\mathbf{A}_{i,j} = NONE$ entails that no arc exists between the nodes. $\mathbf{A}$ is over a label space $\mathbf{L}$: $\mathbf{A} \in \mathbf{L}^{k \times k}$. Each label $l \in \mathbf{L}$ is one of the valid arc labels between two nodes (e.g., *ARG0*, *mod*), in addition to a special label denoting the lack of an arc, *NONE*.

Figure 3: AMR representation of the word *sailor*, which is notable for breaking the word up into a self-contained multi-node unit unpacking the derivational morphology of the word.


Figure 4: AMR representation of the span *January 1, 2008*, an example of how AMR can represent structured data by creating additional nodes like *date-entity* to signify the presence of special structure

2.2 AMR Subgraphs

The mapping from tokens of a sentence to AMR nodes is not one-to-one. A single token or span of tokens can generate a *subgraph* of AMR consisting of multiple nodes. These subgraphs can logically be considered the expression of a single concept, and are useful to treat as such (e.g., see Section 3.1).

Many of these multi-node subgraphs capture structured data such as time expressions (Figure 4). In this example, a *date-entity* node is created to signify that this cluster of nodes is part of a structured sub-component representing a date, where the nodes and arcs within the component have specific semantics. This illustrates a broader recurring pattern in AMR, which is to have an artificial node with certain expected children with special semantics. A particularly salient example of this pattern is the *name* node (see "Rover" in Figure 1) which signifies that all outgoing arcs with label *op* comprise the tokens of a name object.

The ability to decouple the meaning representa-

tion of a lexical item from its surface form allows for rich semantic interpretations of certain concepts in a sentence. For example, the token *sailor* is represented in Figure 3 by a concept graph representing a person who performs the action *sail-01*. Whereas often the AMR node aligned to a span of text is a straightforward function of the text, these cases remain difficult to capture in a more principled way than dictionary lookups.

3 Task decomposition

As in Flanigan et al. (2014), we decompose the end-to-end task of parsing to AMR into two stages: **NER++** – the task of of generating the best AMR sub-graphs given a set of tokens; and **SRL++** – the task of taking the disjoint subgraphs generated by NER++ and creating a fully-connected graph. We describe both tasks in more detail below.

3.1 NER++

Much of the difficulty of parsing to AMR lies in generating local sub-graphs representing the meaning of token spans. For instance, the formalism implicitly demands rich notions of NER, lemmatization, word sense disambiguation, number normalization, and temporal parsing; among others. To illustrate, Figure 2 requires lemmatization (*gleefully* → *glee*), word sense tagging (*run* → *run-01*), and open domain NER (i.e., *Rover*); Figure 4 shows an example of temporal parsing. Furthermore, many of the generated subgraphs (e.g., *sailor* in Figure 3) have rich semantics beyond those produced by standard NLP systems.

We refer to this task as NER++ (see Figure 2). Formally, this is the task of generating a disjoint set of subgraphs representing the meanings of localized spans of words in the sentence.

3.2 SRL++

The second stage of the AMR decomposition consists of generating a coherent graph from the set of disjoint sub-graphs produced by NER++. Unlike the domain-specific arcs within a single sub-graph produced by NER++, the arcs in SRL++ tend to have generally applicable semantics. For example, the SRL arcs (e.g., *ARG0* and *destination* in Figure 2), or the syntactic dependency arcs (e.g., *mod* and *poss* in Figure 2). We use the modified MST algorithm implemented in Flanigan et al. (2014) to produce a labeled DAG corresponding to these

edges.

4 Related Work

The NER++ component of AMR encapsulates a number of lexical NLP tasks. These include named entity recognition (Nadeau and Sekine, 2007; Finkel et al., 2005), word sense disambiguation (Yarowsky, 1995; Banerjee and Pedersen, 2002), lemmatization, and a number of more domain specific tasks. For example, a full understanding of AMR requires normalizing temporal expressions (Verhagen et al., 2010; Strötgen and Gertz, 2010; Chang and Manning, 2012).

In turn, the SRL++ facet of AMR takes many insights from semantic role labeling (Gildea and Jurafsky, 2002; Punyakanok et al., 2004; Srikumar, 2013) to capture the relations between verbs and their arguments. In addition, many of the arcs in AMR have nearly syntactic interpretations (e.g., *mod* for adjective/adverb modification, *op* for compound noun expressions). These are similar to representations used in syntactic dependency parsing (de Marneffe and Manning, 2008; McDonald et al., 2005; Buchholz and Marsi, 2006).

More generally, parsing to a semantic representation is has been explored in depth for when the representation is a logical form (Kate et al., 2005; Zettlemoyer and Collins, 2005; Liang et al., 2011). Recent work has applied semantic parsing techniques to representations beyond lambda calculus expressions. For example, work by Berant et al. (2014) parses text into a formal representation of a biological process. Hosseini et al. (2014) solves algabreic word problems by parsing them into a structured meaning representation. In contrast to these approaches, AMR attempts to capture open domain semantics over arbitrary text.

Interlingua (Mitamura et al., 1991; Carbonell et al., 1999; Levin et al., 1998) are an important inspiration for decoupling the semantics of the AMR language from the surface form of the text being parsed; although, AMR has a self-admitted English bias.

To the best of the authors' knowledge, the JAMR parser (Flanigan et al., 2014) is the only published end-to-end AMR parser. We adopt the dual decomposition based MST algorithm from the parser, and contribute primarily towards replacing the NER++ component.

5 NER++ Method

Our approach to improving NER++ comes from a simple idea: instead of trying to pick which of thousands of AMR sub-graphs to generate from a span of text directly, we partition the AMR sub-graph space in terms of the actions needed to derive a node from its aligned token. At test time we do a sequence labeling of input tokens with these actions, and then deterministically derive the AMR sub-graphs from spans of tokens by applying the transformation decreed by their actions. This dramatically reduces sparsity, and improves generalization (and consequently recall) on the test set, without noticable affecting precision. To illustrate, 38% of the words in the LDC2014E113 test set are unseen during training time.

We describe our partitioning of the action space in Section 5.2, describe our training data in Section 5.4, and our training of the action classifier in Section 5.5.

5.1 Derivation actions

We partition the AMR sub-graph space into a set of 7 actions, each corresponding to an action that will be taken by the NER++ system if a token receives this classification.

- **VERB**: Look for the most similar PropBank frame, make that the title of the corresponding node.

- **IDENTITY**: Take the lowercased version of the token to be the title of the corresponding node.

- **VALUE**: Parse the token to an integer value, and use that as the node. AMR actually does type-check, so

- **LEMMA**: Take the lemma of the token to be the title of the corresponding node.

- **NONE**: Ignore this token in the final output.

- **NAME**: Attach a created "name" node to the top of this span, but don't add an NER action type on top of the "name" node.

- **DICT**: Look up the most probable chunk associate with this lexical span. This functions as a back off if no other actions are appropriate.

5.2 Notes on the DICT action

It's not always possible to derive an AMR subgraph directly from tokens at test time without having memorized a mapping. For example, the parse of "sailor" as "person who sails", see Figure 3, is nearly impossible without some form of memorization. That's where the **DICT** class is important.

To implement a **DICT** class, we memorize a simple mapping from spans of text, like "sailor" to their corresponding most frequently seen AMR sub-graphs in the training data, in this case Figure 3. At test time we can do a lookup in this dictionary for any element that gets labeled with a **DICT** action. Previous approaches have been the equivalent of labeling every node with the **DICT** action, so our reduction of its use is significant. This is the distribution of actions on the LDC2014T12 proxy training data, after our automatic alignment allows us to induce actions (see Section 5.4 for how this is done).

Action	# Tokens	% Total
NONE	41538	0.371
DICT	30027	0.268
IDENTITY	19034	0.170
VERB	11739	0.104
LEMMA	5029	0.045
NAME	4537	0.04
VALUE	16	0.001

Table 1: Distribution of action types in the proxy section of the LDC2014T12 dataset, generated from automatically aligned data.

Note that **DICT** counts for around 27% of the training data, meaning that more than 72% of tokens can be generated correctly by our action type classifier even if we've never seen them before, which is a huge win.

We believe that **DICT** should count for much less than 27%, and **LEMMA** should count for much more than 4%, but issues with existing lemmatizers prevent this, see our error analysis in Section 7.

5.3 Action Informativeness Hierarchy

We define the concept of "action informativeness" of an action a as the probability of deriving the correct node from a span of tokens, given that those tokens are labeled with the action a, and a is the correct action for that span of tokens.
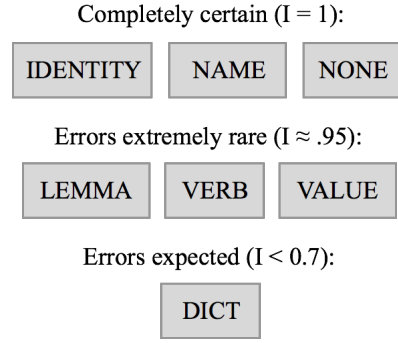

Figure 5: Informativeness hierarchy for action tags within AMR.

To provide a concrete example, our dictionary lookup classifier has a test-set accuracy of 0.67. That means that the "action informativeness" of the **DICT** action is 0.67, because given that we correctly label a token as **DICT**, there is a probability of 0.67 that we correctly generate the corresponding node.

In contrast to **DICT**, correctly labeling a node as **IDENTITY**, **NAME**, and **NONE** have action informativeness of 1.0, since there is no ambiguity in the node generation once one of those actions have been selected, and we are guaranteed (probability 1.0) to generate the correct node given the correct action.

This allows us to induce an action informativeness hierarchy, with more informative actions taking precedence over less informative actions for several important tasks. We demonstrate this hierarchy in Figure 5.

5.4 Inducing Derivation actions from Training Data

Given a set of AMR training data, in the form of (graph,sentence) pairs, we first induce alignments from the graph nodes to the sentence, see Section **??**. Given an alignment, which is an annotation on the graph noting for each node N_i a token S_j that is most likely to have "generated" N_i, we can induce alignments. For concreteness, imagine the token S_j is "running", and the node N_i has the title "run-01". For each action type, we can ask whether that action type is able to take token S_j and correctly generate N_i. The two action types we find that are able to correctly generate this node are **DICT** and **VERB**. We choose the most informative action type of those available to generate the observed node. In this case, that means we choose **VERB**.

In general, our algorithm is as follows. For all S_j to which no N_i exists such that N_i aligns to S_j, assign the action **NONE** to S_j. For all pairs N_i, S_j, assign S_j the most informative action possible that could have generated N_i.

5.5 Action Classifier

We use an extremely simple max-ent classifier to make action decisions. The classifier takes as input a pair $< i, S >$, where i is the index of the token in the input sentence, and S is a sequence tokens representing the source sentence. The output of the classifier is a action T such that the likelihood with respect to the data of token i in sentence S generating a node according to the action specified by T is maximized. See Appendix A for a list of classifier features.

5.6 Test Time Behavior

At test time, given a sequence of input tokens, we do a simple classification of each token separately, to get a sequence labeling of our input tokens. Then for each token, we apply the behavior associated with the token label, and the resulting set of sub-graphs is passed on to SRL++ for linking.

6 Results

6.1 End to end results

Our end to end results are reported by plugging the output of our NER++ into the SRL++ component of JAMR (Flanigan et al., 2014), which is able to produce final AMR graphs when given a sequence of spans and their corresponding chunks. AMR parsing accuracy is measured with a metric called smatch `[citation needed]`, which stands for "s(emantic) match". The metric is the F1 of a best-match between triples implied by the target graph, and triples in the parsed graph. We report much higher recall, and a slightly improved F1 score.

NER++	Dataset	P	R	F1
JAMR	LDC2014T12	**0.671**	0.532	0.59
Robust	LDC2014T12	0.639	**0.596**	**0.62**
JAMR	LDC2014E117	**0.669**	0.529	0.59
Robust	LDC2014E117	0.636	**0.601**	**0.62**

Table 2: Results on two AMR datasets.

6.2 Component results

On our action-type sequence labeling data generated from automatic alignments on train and test splits of *LDC2014T12*, our classifier achieved a test accuracy of **0.841**.

The **DICT** action lookup table achieved an accuracy of **0.67** on the test set. This is remarkable, given that our model moves many of the difficult semantic tasks onto the **DICT** tag, and we are using no learning here beyond a simple count of observed span to sub-graph mappings.

7 Error Analysis

7.1 Weak lemmatization

The **DICT** class was intended to be used for things that a system cannot know without memorization, like the nominalization of "sailor", see Figure 3. These don't occur nearly 25% of the time in the training data. One of the reasons that the **DICT** class is so disappointingly large is that it's stealing from **LEMMA**, because AMR will aggressively normalize words and change their part of speech to a semantic neighbor. For example, 'gleefully' gets mapped to 'glee' and not 'gleeful', which is hard to do automatically with stemming rules in the general case. We leave this as a direction for future work.

7.2 DICT Classifier

The **DICT** class is surprisingly large, and our attempts to handle node generation within the class were fairly feeble.

8 Future Work

8.1 Semantically equivalent POS normalization

The benefit of this approach could be increased by having a very strong stemmer tuned to AMR parsing, which currently doesn't exist.

8.2 Etymological approach to node generation

There is an opportunity to create and test etymologico-semantic approaches to parsing words like 'sailor' that would benefit AMR parsing domain generalization tremendously. We envision a system that upon discovering the unseen noun 'arbitrageur' at training time is able to retrieve the know verb lemma 'arbitrage', and derive that the noun 'arbitrageur' probably refers

to an entity that is engaged in arbitrage. AMR training data provides an opportunity to perform and measure such a task in isolation,

9 Appendix

NER++ Features
Input token
Input token word embedding
Left token
Right token
Left bigram
Right bigram
POS
Left POS
Right POS
Left POS bigram
Right POS bigram
Token's dependency parent token
Token's dependency parent POS
Token's dependency parent arc name
Bag of outgoing dependency arcs
Number of outgoing dependency arcs
Number of outgoing dependency arcs (indicator)
Max JaroWinker to any lemma in PropBank
Closest (JaroWinkler) in PropBank
Token NER
Left NER bigram
Right NER bigram
Right NER bigram
Indicator for if token is a recognized AMR NER type
Indicator for if token is capitalized
Parent arc is prep_* or appos, and parent has NER action
Indicator for token is pronoun
Indicator for token is part of a coref chain
Indicator for token pronoun and part of a coref chain

Table 3: The features for the NER++ maxent classifiers.

References

Laura Banarescu, Claire Bonial, Shu Cai, Madalina Georgescu, Kira Griffitt, Ulf Hermjakob, Kevin Knight, Philipp Koehn, Martha Palmer, and Nathan Schneider. 2013. Abstract meaning representation for sembanking. *Proc. Linguistic Annotation Workshop*.

Satanjeev Banerjee and Ted Pedersen. 2002. An adapted lesk algorithm for word sense disambiguation using wordnet. In *Computational linguistics and intelligent text processing*.

Jonathan Berant, Vivek Srikumar, Pei-Chun Chen, Brad Huang, Christopher D Manning, Abby Vander Linden, Brittany Harding, and Peter Clark. 2014. Modeling biological processes for reading comprehension. In *Proc. EMNLP*.

Sabine Buchholz and Erwin Marsi. 2006. Conll-x shared task on multilingual dependency parsing. In *Proceedings of the Tenth Conference on Computational Natural Language Learning*, pages 149–164. Association for Computational Linguistics.

Jaime G Carbonell, Teruko Mitamura, and Eric H Nyberg. 1999. The kant perspective: A critique of pure transfer (and pure interlingua, pure statistics,...).

Angel Chang and Chris Manning. 2012. SUTIME: a library for recognizing and normalizing time expressions. In *Language Resources and Evaluation*.

Donald Davidson. 1967. The logical form of action sentences. In Nicholas Rescher, editor, *The Logic of Decision and Action*, pages 81–120. University of Pittsburgh Press, Pittsburgh, PA.

Marie-Catherine de Marneffe and Christopher D. Manning. 2008. The Stanford typed dependencies representation. In *Coling 2008: Proceedings of the workshop on Cross-Framework and Cross-Domain Parser Evaluation*.

Jenny Rose Finkel, Trond Grenager, and Christopher Manning. 2005. Incorporating non-local information into information extraction systems by gibbs sampling. In *ACL*.

Jeffrey Flanigan, Sam Thomson, Jaime Carbonell, Chris Dyer, and Noah A Smith. 2014. A discriminative graph-based parser for the abstract meaning representation. In *ACL*.

Daniel Gildea and Daniel Jurafsky. 2002. Automatic labeling of semantic roles. *Computational linguistics*, 28(3):245–288.

Mohammad Javad Hosseini, Hannaneh Hajishirzi, Oren Etzioni, and Nate Kushman. 2014. Learning to solve arithmetic word problems with verb categorization. In *EMNLP*.

Rohit J. Kate, Yuk Wah Wong, and Raymond J. Mooney. 2005. Learning to transform natural to formal languages. In *AAAI*, Pittsburgh, PA.

Lori S Levin, Donna Gates, Alon Lavie, and Alex Waibel. 1998. An interlingua based on domain actions for machine translation of task-oriented dialogues. In *ICSLP*, volume 98, pages 1155–1158.

P. Liang, M. I. Jordan, and D. Klein. 2011. Learning dependency-based compositional semantics. In *ACL*.

Ryan McDonald, Koby Crammer, and Fernando Pereira. 2005. Online large-margin training of dependency parsers. In *ACL*, Morristown, NJ, USA.

Teruko Mitamura, Eric H Nyberg, and Jaime G Carbonell. 1991. An efficient interlingua translation system for multi-lingual document production.

David Nadeau and Satoshi Sekine. 2007. A survey of named entity recognition and classification. *Lingvisticae Investigationes*, 30(1):3–26.

Terence Parsons. 1990. *Events in the Semantics of English: A study in subatomic semantics*. MIT Press, Cambridge, MA.

Vasin Punyakanok, Dan Roth, Wen-tau Yih, and Dav Zimak. 2004. Semantic role labeling via integer linear programming inference. In *Proceedings of the 20th international conference on Computational Linguistics*, page 1346. Association for Computational Linguistics.

Vivek Srikumar. 2013. *The semantics of role labeling*. Ph.D. thesis, University of Illinois at Urbana-Champaign.

Jannik Strötgen and Michael Gertz. 2010. Heideltime: High quality rule-based extraction and normalization of temporal expressions. In *Proceedings of the 5th International Workshop on Semantic Evaluation*, Sem-Eval.

Marc Verhagen, Roser Sauri, Tommaso Caselli, and James Pustejovsky. 2010. Semeval-2010 task 13: TempEval-2. In *Proceedings of the 5th International Workshop on Semantic Evaluation*, Uppsala, Sweden.

David Yarowsky. 1995. Unsupervised word sense disambiguation rivaling supervised methods. In *ACL*.

Luke S. Zettlemoyer and Michael Collins. 2005. Learning to map sentences to logical form: Structured classification with probabilistic categorial grammars. In *UAI*. AUAI Press.